

# **Speculations on Test-Time Scaling**

Sasha Rush Daniel Ritter

Cornell

# Outline

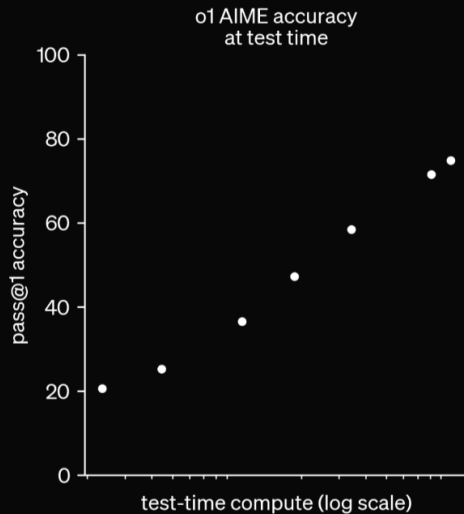

Introduction

The Clues

Notation

The Suspects

Conclusions



# Context

- LLM (2018-2024) driven by training scaling
- Speculation: Benefit of static data running out

# Implication

- Breakthrough in large-scale RL Training
-

# What have we seen?

- Public demo model
- Strong result in constrained domains.

# This Talk

- Survey of the public literature
- Synthesis of discussions with expert
- Gossip and hearsay

# Thanks

Lewis Tunstall, Edward Beeching, Aviral Kumar, Charlie Snell,  
Michael Hassid, Yoav Artzi, Risab Agarwal, Kanishk Gandhi,  
Wenting Zhao, Yuntian Deng, Nathan Lambert



# What we know

Our large-scale **reinforcement learning algorithm** teaches the model how to think productively using its **chain of thought** in a highly **data-efficient** training process.

# What we know

- RL; Signal from verifiable problems
- CoT; “Thinking” occurs in token stream
- Data Efficient; Fixed set of good problems

# From Gossip

- Single final model
- Not learned from expert examples
-

# Chain of Thought

o1 learns to hone its chain of thought and refine the strategies it uses. It learns to recognize and **correct its mistakes**. It learns to **break down tricky steps** into simpler ones. It learns to try a **different approach** when the current one isn't working.

# Review: Chain of Thought



# Planning

*First, the cipher seems connected to the plaintext.  
Given the time constraints, perhaps the easiest way  
is to try to see patterns.*

*Option 1: Try to find mapping from letters to letters.  
Do any letters match?*

*First, let's write down the ciphertext and plaintext letters on top of each other.*

# Backtracking

# Strategies



# Summary

- Solves problems by very long CoT
- CoT includes “thinking” (search / planning)
- Core novelty: Inducing this behavior

# Notation: LLM Sampling (No learning yet!)


- $x$ ; problem specification
- $z_{1:T} \in \mathcal{S}^T$ ; chain of thought (CoT) steps
- $y \in \mathcal{Y}$ ; final answer

$$p(y|x) = \mathbb{E}_z p(y|x, z)$$

# Warm-up: Ancestral Sampling

$$z_{1:T} \sim p(\cdot | x)$$

$$y \sim p(\cdot | x, z_{1:T})$$



$T$  is the amount of test-time compute


# Warm-up: Monte-Carlo (Self-Consistency)

For  $N$  samples,

$$z_{1:T} \sim p(\cdot | \mathbf{x})$$

$$y^n \sim p(\cdot | \mathbf{x}, z_{1:T})$$

Pick majority choice  $y^n$



# Assumption: Verifier

$$\text{Ver}_x : \mathcal{Y} \rightarrow \{0, 1\}$$

Examples:

- Regular expression for math
- Unit test for code
- Test questions for science


# Warm up: Rejection Sampling / Best-of-N

For  $n = 1$  to  $N$ :

$$z^n \sim p(z|x)$$

$$y^n \sim p(y|x, z^n)$$

Verified set  $\{y^n : \text{Ver}(y^n)\}$




# Warm up: Monte-Carlo Roll-Outs

Given partial CoT  $z_{1:t}$ , expected value,

$$\mathbb{E}_{y \sim p(y|z, x), z_{t:T}} \text{Ver}(y)$$

Rollout = Monte Carlo for this expectation.



# Goal: Learning with Latent CoTs

Maximum likelihood;

$$\max_{\theta} \sum \log p(y|x; \theta) = \\ \sum \log \mathbb{E}_z p(y|x, z; \theta)$$

Classic combinatorial  
expectation



# Outline

Introduction

The Clues

Notation

The Suspects


Conclusions

# The Suspects

- Guess + Check
- Guided Search
- AlphaZero
- Learn to Search
- Wildcard

# Suspect 1: Guess + Check

- 1) Sample  $N$  CoTs
- 2) Check if successful
- 3) Train on good ones



# G+C Formalization: Rejection Sampling EM

$$\max_{\theta} \sum \log E_{z \sim p(z|x; \theta)} p(y|x, z)$$

- E-Step: For  $n = 1$  to  $N$ :

$$z^n \sim p(\cdot|x)$$

$$y^n \sim p(\cdot|x, z^n)$$

Keep verified set  $\mathcal{Z} = \{z^n : \text{Ver}(y^n)\}$

- M-Step: Fit  $\theta' \leftarrow \arg \max_{\theta} \sum_{z \in \mathcal{Z}} \log p(z|x; \theta)$

# G+C Variants

- STaR
- ReST
- ReST-EM
- Filtered Rejection Sampling
- Best-of-N Training

## G+C Variants

- Batched -> Compute trajectories first, then train with behavioral cloning
- Online -> Use policy gradient-like steps to update after each example

# G+C Empirical Results




Figure 5 | **Pass@K results** for PaLM-2-L pretrained model as well as model fine-tuned with ReST<sup>EM</sup>. For a fixed number of samples K, fine-tuning with ReST<sup>EM</sup> substantially improves Pass@K performance. We set temperature to 1.0 and use nucleus sampling with  $p = 0.95$ .

## G+C Why might this be right?

- Extremely simple and scalable
- Good baseline in past work
- No evidence this learns to correct, plan
- Well-explored in literature with marginal gains



# More Structure?

- Rejection sampling may be really inefficient.
- Particularly on hard problems, may get no signal

## Suspect 2: Guided Search

- During CoT sampling, use a heuristic to improve trajectories
- Check if final versions are successful
- Train on good ones

# GS: Beam Search with Guide


$r : \mathcal{S}^t \rightarrow \mathbb{R}$ ; Guide function

For each step  $t$

1. Sample next step,

$$z_t \sim p(\cdot | \mathbf{x}, z_{1:t-1}^i)$$

2. Keep the top  $N$  samples,  
ordered by  $r(z_t)$



# What to use as Guide?


- Monte Carlo Roll-outs
- Learned Value Function

# Beam Search with Roll-Outs

For a  $z_t$ , sample answers


$$y^n \sim p(\cdot | x, z_{1:t-1})$$

$$r_{MC}(z_t) = \frac{1}{N} \sum_{i=1}^N \text{Ver}(y^i)$$



# Amortized Roll-Outs

- Rollouts are costly, so instead learn a model  $r_\psi(z_t)$  to approximate rollouts
- Use  $r_{MC}$  to determine labels to train  $r_\psi$



# What about test time?

- Learned rewards can improve test-time without verifier.
-

# Terminology

[Uesato et al., 2022, Setlur et al., 2024,  
Wang et al., 2023, Lightman et al., 2023,  
Snell et al., 2024]

- Value
- PRM
- PAV
- Math Shepard.
- snell.



# Why might this be right?


- OpenAI is exploring
- Makes RS more efficient.
- Learned rewards are effective
- Assumption: o1 is a single test-time model (although could train or distill-in)
- Not clear if it learns planning.

# More Structure


- Improving search seems critical.

# Reminder: AlphaZero

**a** Self-play



**b** Neural network training



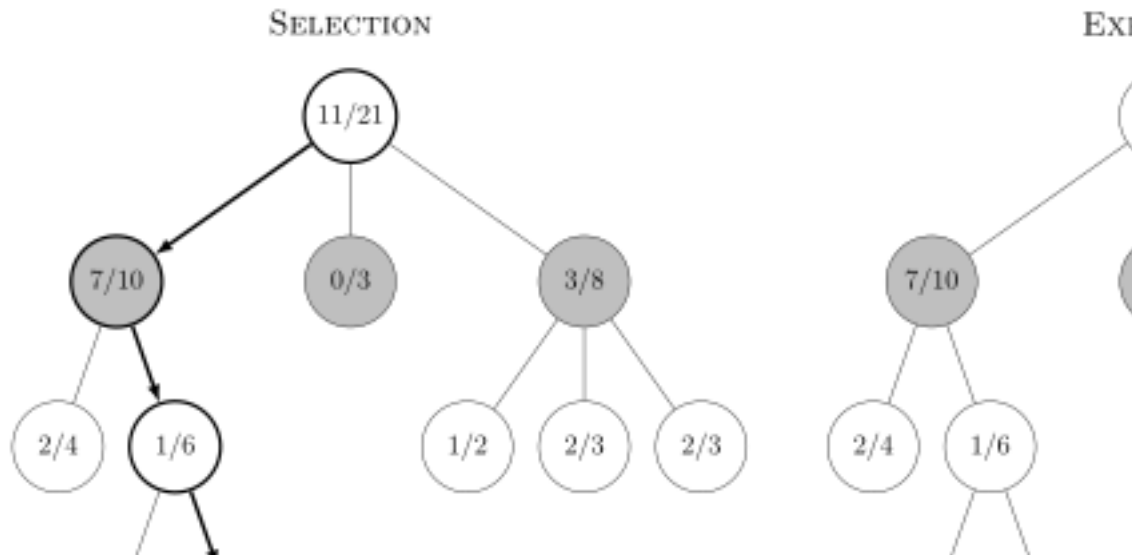
## Suspect 3: AlphaZero

- Self-play using guided-search with exploration
- Label final outcomes of self-play games
- Train guide and generator

# Formalized: Expert Iteration

- Iterative algorithm combining learned model + expert search with a verifier.
- Generate samples using  $p(y, z|x)$ , reward model  $r(z_t)$ , and search algorithm (e.g. beam search)
- Label samples using  $Ver(y)$
- Train  $p(y, z|x)$ ,  $r(z_t)$  on the labeled samples, and repeat

# MCTS exploration



# UCB for Language

- Selection: Walk down tree to leaf  $z_{t-1}$
- Expand: Sample  $K$  next steps  $z_t^i$ , pick one at random
- Rollouts: Sample  $z_{t+1} \dots z_T$
- Backprop: Update nodes counts  $z_{1:t}$  based on results

# Compared with Search

- System builds in exploration
- Scales to more train-time search
- Costly to maintain open states
- More complex algorithmically



# Empirical Results




Figure 8: Test accuracy of 13B V-STaR compared to baselines. We report Best-of-64 for verification-based methods and Pass@1 for others. **(Left)** Test accuracy for training tasks. **(Right)** Transfer evaluation of GSM8K and MBPP trained models on MATH subset and HumanEval respectively.

# Why might this be right?

- Major demonstrated RL result
- 
- 
-

# More Structure

- Can we force the model to search?

## Suspect 4: Learning to Correct

- Sample  $N$  Successful CoTs
- Edit  $z \rightarrow z'$  to inject incorrect expansions before correct ones.
- Train on  $z'$  trajectories

# Formalized: Stream of Search



# Empirical Results

Score results

## Why might this be right?

- 

-

## Why might this be wrong?

- 

-



# Less Structure?

- Maybe this is all too much...

# Suspect 5: Compute Injection

-



# Reference I

- [Anthony et al., 2017] Anthony, T., Tian, Z., and Barber, D. (2017).  
Thinking fast and slow with deep learning and tree search.  
*arXiv [cs.AI]*.
- [Brown et al., 2024] Brown, B., Juravsky, J., Ehrlich, R., Clark, R., Le, Q. V., Ré, C., and Mirhoseini, A. (2024).  
Large language monkeys: Scaling inference compute with repeated sampling.  
*arXiv [cs.LG]*.

## Reference II

[Gandhi et al., 2024] Gandhi, K., Lee, D., Grand, G., Liu, M., Cheng, W., Sharma, A., and Goodman, N. D. (2024). Stream of search (SoS): Learning to search in language. *arXiv [cs.LG]*.

[Gulcehre et al., 2023] Gulcehre, C., Paine, T. L., Srinivasan, S., Konyushkova, K., Weerts, L., Sharma, A., Siddhant, A., Ahern, A., Wang, M., Gu, C., Macherey, W., Doucet, A., Firat, O., and de Freitas, N. (2023). Reinforced self-training (ReST) for language modeling. *arXiv [cs.CL]*.

## Reference III

[Lightman et al., 2023] Lightman, H., Kosaraju, V., Burda, Y., Edwards, H., Baker, B., Lee, T., Leike, J., Schulman, J., Sutskever, I., and Cobbe, K. (2023).

Let's verify step by step.

*arXiv [cs.LG].*

[Setlur et al., 2024] Setlur, A., Nagpal, C., Fisch, A., Geng, X., Eisenstein, J., Agarwal, R., Agarwal, A., Berant, J., and Kumar, A. (2024).

Rewarding progress: Scaling automated process verifiers for LLM reasoning.

*arXiv [cs.LG].*

## Reference IV

[Singh et al., 2023] Singh, A., Co-Reyes, J. D., Agarwal, R., Anand, A., Patil, P., Garcia, X., Liu, P. J., Harrison, J., Lee, J., Xu, K., Parisi, A., Kumar, A., Alemi, A., Rizkowsky, A., Nova, A., Adlam, B., Bohnet, B., Elsayed, G., Sedghi, H., Mordatch, I., Simpson, I., Gur, I., Snoek, J., Pennington, J., Hron, J., Kenealy, K., Swersky, K., Mahajan, K., Culp, L., Xiao, L., Bileschi, M. L., Constant, N., Novak, R., Liu, R., Warkentin, T., Qian, Y., Bansal, Y., Dyer, E., Neyshabur, B., Sohl-Dickstein, J., and Fiedel, N. (2023).

Beyond human data: Scaling self-training for problem-solving with language models.

## Reference V

*arXiv [cs.LG].*

[Snell et al., 2024] Snell, C., Lee, J., Xu, K., and Kumar, A. (2024).

Scaling LLM test-time compute optimally can be more effective than scaling model parameters.

*arXiv [cs.LG].*

[Uesato et al., 2022] Uesato, J., Kushman, N., Kumar, R., Song, F., Siegel, N., Wang, L., Creswell, A., Irving, G., and Higgins, I. (2022).

Solving math word problems with process- and outcome-based feedback.



## Reference VI

*arXiv [cs.LG].*

[Wang et al., 2023] Wang, P., Li, L., Shao, Z., Xu, R. X., Dai, D., Li, Y., Chen, D., Wu, Y., and Sui, Z. (2023).

Math-shepherd: Verify and reinforce LLMs step-by-step without human annotations.

*arXiv [cs.AI].*

[Zelikman et al., 2022] Zelikman, E., Wu, Y., Mu, J., and Goodman, N. D. (2022).

STaR: Bootstrapping reasoning with reasoning.

*arXiv [cs.LG].*